

Question #8

Topic 4

You plan to store data in Azure Blob storage for many years. The stored data will be accessed rarely.

You need to ensure that the data in Blob storage is always available for immediate access. The solution must minimize storage costs.

Which storage tier should you use?

A. Cool

B. Archive

C. Hot

**Correct Answer:** A

Azure cool tier is equivalent to the Amazon S3 Infrequent Access (S3-IA) storage in AWS that provides a low cost high performance storage for infrequently access data.

Note: Azure's cool storage tier, also known as Azure cool Blob storage, is for infrequently-accessed data that needs to be stored for a minimum of 30 days.

Typical use cases include backing up data before tiering to archival systems, legal data, media files, system audit information, datasets used for big data analysis and more.

The storage cost for this Azure cold storage tier is lower than that of hot storage tier. Since it is expected that the data stored in this tier will be accessed less frequently, the data access charges are high when compared to hot tier. There are no additional changes required in your applications as these tiers can be accessed using APIs in the same manner that you access Azure storage.

Incorrect Answers:

B: Even though Azure archive storage offers the lowest cost in terms of data storage, its data retrieval charges are higher than that of hot and cool tiers. In fact, the data in the archive tier remains offline until the tier of the data is changed using a process called hydration. The process of hydrating data in the archive storage tier and moving it to either hot or cool tier could take up to 15 hours and, hence, it is only intended for data that can afford that kind of access delay.

C: The storage cost for this Azure cold storage tier is lower than that of hot storage tier.

References:

<https://cloud.netapp.com/blog/low-cost-storage-options-on-azure>

Question #9

Topic 4

**HOTSPOT -**

You have a virtual machine scale set named SS1.

You configure autoscaling as shown in the following exhibit.

Delete warning The very last or default recurrence rule cannot be deleted. Instead, you can disable autoscale to turn off autoscale.

Scale mode ☒ Scale based on a metric ☐ Scale to a specific instance count

Scale out

When	SS1	(Average) Percentage CPU > 75	Increase instance count by 3
------	-----	-------------------------------	------------------------------

Rules

Scale in

When	SS1	(Average) Percentage CPU < 25	Decrease instance count by 2
------	-----	-------------------------------	------------------------------

+ Add a rule

Instance limits

Minimum ⓘ

Maximum ⓘ

Default ⓘ

Schedule This scale condition is executed when none of the other scale condition(s) match

You configure the scale out and scale in rules to have a duration of 10 minutes and a cool down time of 10 minutes.  
Use the drop-down menus to select the answer choice that answers each question based on the information presented in the graphic.  
NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

If SS1 scales to nine virtual machines, what is the minimum amount of time before SS1 will scale up?

▼

10 minutes

20 minutes

30 minutes

60 minutes

If SS1 scales to nine virtual machines, and then the average processor utilization is 30 percent for one hour, how many virtual machines will be in SS1?

▼

1

3

6

9

12

15

Answer Area

If SS1 scales to nine virtual machines, what is the minimum amount of time before SS1 will scale up?

▼

10 minutes

20 minutes

30 minutes

60 minutes

Correct Answer:

If SS1 scales to nine virtual machines, and then the average processor utilization is 30 percent for one hour, how many virtual machines will be in SS1?

▼

1

3

6

9

12

15

Box 1: 20 Minutes. 10 minutes cool down time after the last scale-up plus 10 minutes duration equals 20 minutes.  
Box 2: 9 virtual machines. 30% does not match the scale in requirement of less than 25% so the number of virtual machines will not change.

HOTSPOT -

You have 20 Azure virtual machines that run Windows Server 2016 based on a custom virtual machine image. Each virtual machine hosts an instance of a VSS- capable web app that was developed in-house. Each instance is accessed by using a public endpoint. Each instance uses a separate database. The average database size is 200 GB.

You need to design a disaster recovery solution for individual instances. The solution must meet the following requirements:

- ☞ Provide a recovery time objective (RTO) of six hours
- ☞ Provide a recovery point objective (RPO) of eight hours
- ☞ Support recovery to a different Azure region
- ☞ Support VSS-based backups
- ☞ Minimize costs

What should you include in the recommendation? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Disaster recovery solution:

	▼
Azure Site Recovery	
Scheduled backups by using an Azure Backup agent	
Scheduled backups by using Windows Server Backup	
Scheduled virtual machine-level backups	

Storage type:

	▼
A Recovery Services vault	
Premium managed disks	
Unmanaged disks in an storage account that uses RA-GRS	

Correct Answer:

Answer Area

Disaster recovery solution:

	▼
Azure Site Recovery	
Scheduled backups by using an Azure Backup agent	
Scheduled backups by using Windows Server Backup	
Scheduled virtual machine-level backups	

Storage type:

	▼
A Recovery Services vault	
Premium managed disks	
Unmanaged disks in an storage account that uses RA-GRS	

HOTSPOT -

You plan to deploy the backup policy shown in the following exhibit.



Policy1

Associated items

Delete

Save

Discard

Backup frequency

Daily

6:00 PM

(UTC) Coordinated Universal Time

Retention range

Retention of daily backup point.

At

For

6:00 PM

90

Day(s)

Retention of weekly backup point.

On

At

For

Sunday

6:00 PM

26

Week(s)

Retention of monthly backup point.

Week Based

Day Based

On

Day

At

For

First

Sunday

6:00 PM

36

Month(s)

Retention of yearly backup point.

Not Configured

Use the drop-down menus to select the answer choice that completes each statement based on the information presented in the graphic.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Virtual machines that are backed up using the policy can be recovered for up to a maximum of [answer choice].

▼

90 days

26 weeks

36 months

45 months

The minimum recovery point objective (RPO) for virtual machines that are backed up by using the policy is [answer choice].

▼

1 hour

1 day

1 week

1 month

1 year

Correct Answer:

Answer Area

Virtual machines that are backed up using the policy can be recovered for up to a maximum of [answer choice].

▼

90 days

26 weeks

36 months

45 months

The minimum recovery point objective (RPO) for virtual machines that are backed up by using the policy is [answer choice].

▼

1 hour

1 day

1 week

1 month

^

HOTSPOT -

You have databases in Azure as shown in the following table.

Name	Type	Pricing tier
SQLdb1	Azure SQL Database	Basic
SQLdb2	Azure SQL Database	Standard
SQLdb3	Azure SQL Database	Premium
DW1	Azure SQL Data Warehouse	Not available

You are designing a data retention policy.

You need to identify which databases can retain a daily backup for up to 35 days and which databases can retain monthly backups for up to 120 months.

Which databases should you identify? To answer, select the appropriate options in the answer area.

NOTE: Each correct selection is worth one point.

Hot Area:

Answer Area

Databases that can retain a daily backup for up to 35 days:

DW1 only

SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2, SQLdb3 and DW1

Databases that can retain monthly backups for up to 120 months:

DW1 only

SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2, SQLdb3 and DW1

Answer Area

Correct Answer:

Databases that can retain a daily backup for up to 35 days:

DW1 only

SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2, SQLdb3 and DW1

Databases that can retain monthly backups for up to 120 months:

DW1 only

SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2 and SQLdb3 only

SQLdb1, SQLdb2, SQLdb3 and DW1